

Interrupted Time Series & Regression Discontinuity

Control Units



None

Some

Time-Points



Two

Pre-Post

Diff-in-Diff

Many

Interrupted
Time-Series
(ITS)

Synthetic
Control

	Two	Pre-Post	Diff-in-Diff
	Many	Interrupted Time-Series (ITS)	Synthetic Control

The story so far

The **proposition 99** data has a number of pre- and post-intervention observations (i.e. time points)

So far we **computed averages** and estimated

$$\overline{CE}_{post} = \bar{Y}_{post}^1 - \bar{Y}_{post}^0$$

Interrupted Time Series

Interrupted Time Series:

- Instead of taking averages, use pre-intervention data Y_{pre}^0 to **forecast/predict** Y_{post}^0
- Once we have predictions \hat{Y}_{post}^0 , we compare those to the observed Y_{post}^1
- I.e. we use pre-intervention data to **impute** the missing counterfactual

This means we can in principle estimate

$$\widehat{CE}_t = Y_t^1 - \widehat{Y}_t^0$$

<i>Time</i>	Y_t	A_t	Y_t^0	Y_t^1
1	7	0	7	NA
2	9	0	9	NA
3	6	0	6	NA
4	5	0	5	NA
5	6	0	6	NA
6	2	1	NA	2
7	3	1	NA	3
8	1	1	NA	1
...
T	2	1	NA	2

Interrupted Time Series

<i>Time</i>	Y_t	A_t	Y_t^0	Y_t^1
1	7	0	7	NA
2	9	0	9	NA
3	6	0	6	NA
4	5	0	5	NA
5	6	0	6	NA
6	2	1	NA	2
7	3	1	NA	3
8	1	1	NA	1
...
<i>T</i>	2	1	NA	2

Fit a forecasting Model
 $\hat{Y}_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \beta * Time$

Interrupted Time Series

<i>Time</i>	Y_t	A_t	Y_t^0	Y_t^1
1	7	0	7	NA
2	9	0	9	NA
3	6	0	6	NA
4	5	0	5	NA
5	6	0	6	NA
<hr/>				
6	2	1	\widehat{Y}_6^0	2
7	3	1	\widehat{Y}_7^0	3
8	1	1	\widehat{Y}_8^0	1
...
T	2	1	\widehat{Y}_T^0	2

Fit a forecasting Model
 $\widehat{Y}_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \beta * Time$

Make forecasts

Interrupted Time Series

<i>Time</i>	Y_t	A_t	Y_t^0	Y_t^1
1	7	0	7	NA
2	9	0	9	NA
3	6	0	6	NA
4	5	0	5	NA
5	6	0	6	NA
6	2	1	\widehat{Y}_6^0	2
7	3	1	\widehat{Y}_7^0	3
8	1	1	\widehat{Y}_8^0	1
...
T	2	1	\widehat{Y}_T^0	2

Fit a forecasting Model
 $\widehat{Y}_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \beta * Time$

Make forecasts

$$\widehat{CE}_t = Y_t^1 - \widehat{Y}_t^0$$

Interrupted Time Series

Point forecasts allow us to compute point estimates of our causal effect

$$\widehat{CE}_t = Y_t^1 - \widehat{Y}_t^0$$

We can quantify our **uncertainty** about the causal effect based on our **uncertainty** around our (model-based) forecasts

Building a forecasting model

Much of the challenge of this approach is in choosing an appropriate **forecasting** model

These can be very simple or very complex, e.g.:

- If we forecast with the **mean** we are very close to the post – pre analysis

$$Y_t = \mu_{pre} + e_t$$

Building a forecasting model

Much of the challenge of this approach is in choosing an appropriate **forecasting** model

These can be very simple or very complex, e.g.:

- If we forecast with the **mean** we are very close to the post – pre analysis

$$Y_t = \mu_{pre} + e_t$$


- If we forecast with a **growth curve** we would model the overall time trend

$$Y_t = \beta_0 + \beta_1 Time + e_t$$

Forecasting with growth curves

```
# predict pre-intervention sales by year
fit_growth <- lm(
  formula = cigsale ~ year,
  prop99_ts > filter(prepost == "Pre")
)

# predict values for the post-intervention period
pred <- predict(
  object = fit_growth,
  newdata = prop99_ts,
  interval = "prediction"
)
```



Building a forecasting model

Much of the challenge of this approach is in choosing an appropriate **forecasting** model

These can be very simple or very complex, e.g.:

- If we forecast with the **mean** we are very close to the post – pre analysis

$$Y_t = \mu_{pre} + e_t$$

- If we forecast with a **growth curve** we would model the overall time trend

$$Y_t = \beta_0 + \beta_1 Time + e_t$$

- We can also forecast by using **time-series models** that model **autocorrelation**

$$Y_t = \phi_1 Y_{t-1} + e_t \quad Y_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + e_t \quad Y_t - Y_{t-1} = \gamma e_{t-1} + e_t$$

e.g. ARIMA models can account for autocorrelation and time trends


Fitting time-series models fpp3

Prep data as tsibble object

```
prop99_ts <-  
  prop99 ▷  
  filter(state = "California") ▷  
  select(year, cigsale) ▷  
  mutate(prepost = factor(year > 1988, labels = c("Pre", "Post"))) ▷  
  as_tsibble(index = year) ▷  
  mutate(year0 = year - 1989)
```

Fit a model (automatic model selection on AR, differencing, and MA components)
and produce forecasts

```
fit_arima <-  
  prop99_ts ▷  
  filter(prepost = "Pre") ▷  
  model(timeseries = ARIMA(cigsale, ic = "aicc"))  
  
# create forecasts  
fcsts <- forecast(fit_arima, h = "12 years")
```



Key Assumptions

Our inferences about the causal effect are entirely dependent on being able to fit **an appropriate forecasting model**

- i.e. one that correctly captures the trend and autocorrelation structures in the data

In practice, this may be very difficult

Key Assumptions

Data driven approaches can be applied, but may only be feasible with **a large amount of pre-intervention training data**

- We use information criteria for model selection
- See also: cross-validation

In addition, different forecasting models come with their own assumptions,

- E.g. **constant trend** or **time-invariant relationships**

Poor forecasts = Poor estimates (and uncertainty) of causal effects

Key Assumptions

When comparing to the **pre-post design**;

- We relax the no-trend assumption: we model any trend / serial dependence

No-confounding assumption:

- We still assume that any changes can be attributed to the intervention
- And not, e.g., something else that happened around the same time
- To tackle that we need control units + other assumptions

Regression Discontinuity (RDD)

Closely related technique, but used in many other contexts

E.g., instead of “Time” we may have “Income”; if above X, eligible for social welfare.

In a RDD analysis you fit **piecewise growth-curve** type model such as

$$Y_t = \beta_0 + \beta_1 A_t + \beta_2 \text{ Time} + \beta_3 * \text{Time} * A_t + e_t$$

In this model the effect of the intervention is parameterized by the change in **level** β_1 and the change in **trend** β_3 after the intervention

Hypothesis tests on these parameters are used as hypothesis tests about the presence / absence of a causal effect

Regression Discontinuity in Practice

```
fit_rdd <- lm(cigsale ~ year0 + prepost + year0:prepost, prop99_ts)
summary(fit_rdd)
```

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	98.4158	2.4746	39.770	< 2e-16	***
year0	-1.7795	0.2170	-8.199	8.36e-09	***
prepostPost	-20.0581	3.7471	-5.353	1.18e-05	***
year0:prepostPost	-1.4947	0.4846	-3.084	0.00467	**

signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.182 on 27 degrees of freedom

Multiple R-squared: 0.9732, Adjusted R-squared: 0.9702

F-statistic: 326.4 on 3 and 27 DF, p-value: < 2.2e-16



Regression Discontinuity

Basic Idea:

You directly **model** whatever changes you think happen to the target process
- Instead of making forecasts/predictions of the counterfactual directly

Advantages

- More direct. Inference about CE based on significance tests on “change” parameters
- Many extensions and theory to deal with, e.g., “sharp” vs “fuzzy” designs

Disadvantages

- Strongly rely on correct model specification and model interpretability; specify “where” or “how” the intervention has an effect

Practical

Work in your groups!

Lunch